

ROHAN KRISHNAN

Krishnan.Rohan@outlook.com | (717) 839-1010 | Philadelphia, PA | github.com/r0hankrishnan | linkedin.com/in/rohankrish

EDUCATION

M.S.E. in Data Science, University of Pennsylvania, Philadelphia, PA Expected 2027

- Relevant Coursework: Big Data Analytics, Applied Machine Learning, Statistics for Data Science

Hons. B.S. in Economics & B.S.B.A. in Business Information Systems, University of Pittsburgh, Pittsburgh, PA 2020 – 2024

- Minor in Applied Statistics & Certificate in Business Analytics
- Honors/Awards: Summa Cum Laude, \$40,000 Merit Scholarship, \$4,000 Honors Research Fellowship, \$3000 DSG Analytics Case Competition 1st Prize, PLUS 3 International Business Fellow, Haller Global Honors Fellow

WORK EXPERIENCE

Team Lead, Paragon Technology Policy Fellowship, Fort Collins, CO Jun 2026 – Aug 2026

- Lead a 7-person team assessing AI suitability for the City of Fort Collins' digital services and acted as primary contact for City IT and Communications leadership.
- Ran multiple resident focus group interviews and conducted desk research on census and labor demographics across 7 city services and 20+ municipal AI deployments to map the access barriers residents face.
- Delivered two stakeholder assessments — resident service access and AI suitability — each with an executive one-pager scoping where AI fits, where it fails, and how to evaluate it before building.

Business Analyst, Penn Graduate Consulting Club / AI Maternal Health Startup, Philadelphia, PA Mar 2026 – May 2026

- Audited and rebuilt a payer value-based pricing model in Excel for an AI early-pregnancy platform, quantifying \$255K–\$1.44M in annual payer savings and a defensible \$19–\$94 per-pregnancy price point across low/base/high scenarios for key payer clients in the Philadelphia region.
- Corrected a ~7x market-sizing error by deriving Philadelphia's privately insured pregnancy volume from CDC WONDER natality data and PA Dept. of Health coverage reports.
- Re-sourced all 15+ model assumptions from founder estimates to primary evidence — peer-reviewed studies, CPT-level cost benchmarks, and payer data— reframing the value narrative around two clinically defensible drivers, ED diversion and care-setting optimization.

Freelance Data & Automation Consultant, Self-Employed / Local Healthcare Clinic, Mechanicsburg, PA Jul 2024 – Nov 2024

- Built a HIPAA-compliant Python ETL pipeline against a \$400K billing backlog, cutting manual data-prep time 70%.
- Prototyped an audio-to-structured-data transcription tool for physician notes, with testing indicating that the tool cut duplicate data entry 90% and would shorten revenue-cycle resolution by 45+ days.

Technology Co-Lead, Pitt Business Advanced Analytics Consulting / Metal Manufacturer, Pittsburgh, PA Jan 2023 – Apr 2023

- Co-led a 12-person team building an ensemble price-forecasting model (ARIMA, Prophet, random forest) for a national metals manufacturer, reaching 85% F1 on price-movement classification and MAE within \$3 on spot price.
- Shipped an R Shiny dashboard surfacing model predictions and procurement metrics to guide \$5M+ in annual zinc purchasing decisions.

PROJECTS

Frame Finder, Hybrid semantic search engine for tennis racquet discovery May 2026 – Jun 2026

- Scraped and cleaned a 333-racquet dataset from 11 brand catalogs, engineering 15+ structured features from inconsistent free-text specs.
- Distilled raw marketing copy into standardized descriptions via LLM before embedding, improving match quality on feel-based queries.
- Combined BM25 with dense embeddings via reciprocal rank fusion, raising Hit@5 15 points over keyword-only search across both playstyle and exact-model lookups.
- Deployed as a live FastAPI service logging real user searches and feedback to benchmark retrieval quality.

Flight Delay Triage System, ML-based operational risk ranking for incoming PHL flights Feb 2026 – Mar 2026

- Ranked inbound flights by arrival-delay probability on merged BTS and Open-Meteo weather data, hitting 96% precision at k of 500 and 100% through k of 100 compared to a 19% base rate.
- Selected XGBoost over logistic regression, random forest, and tree baselines, evaluating via average precision rather than on accuracy due to class imbalance.
- Deployed as a live Streamlit dashboard and FastAPI scoring endpoint for operational triage.

SKILLS

Python (pandas, scikit-learn, NumPy, FastAPI, matplotlib, BeautifulSoup4, Streamlit, PySpark), R, SQL, Tableau, Git, Ollama